



RAPIDS

# Reconciling Availability, Accuracy, and Performance in Managing Geo-Distributed Scientific Data

Lipeng Wan, Jieyang Chen, Xin Liang, Ana Gainaru, and others

Muhammad Raees  
mr2714@rit.edu

# Objective

Increase the performance of geo-distributed data (scientific) storage systems by

- Reducing storage and network overhead
- Improving data access/latency

While maintaining the same level of data

- Availability
- Accuracy


# Data Size Grows

- Large scientific datasets
  - XGC gyrokinetic particle-in-cell code
  - Square Kilometer Array (SKA)
- High-Performance Storage Needs
- Managing failures is critical
  - Availability, accuracy, and performance

# Improving Resilience

## Data Duplication

- Availability Improvement
- Storage and Network Overhead



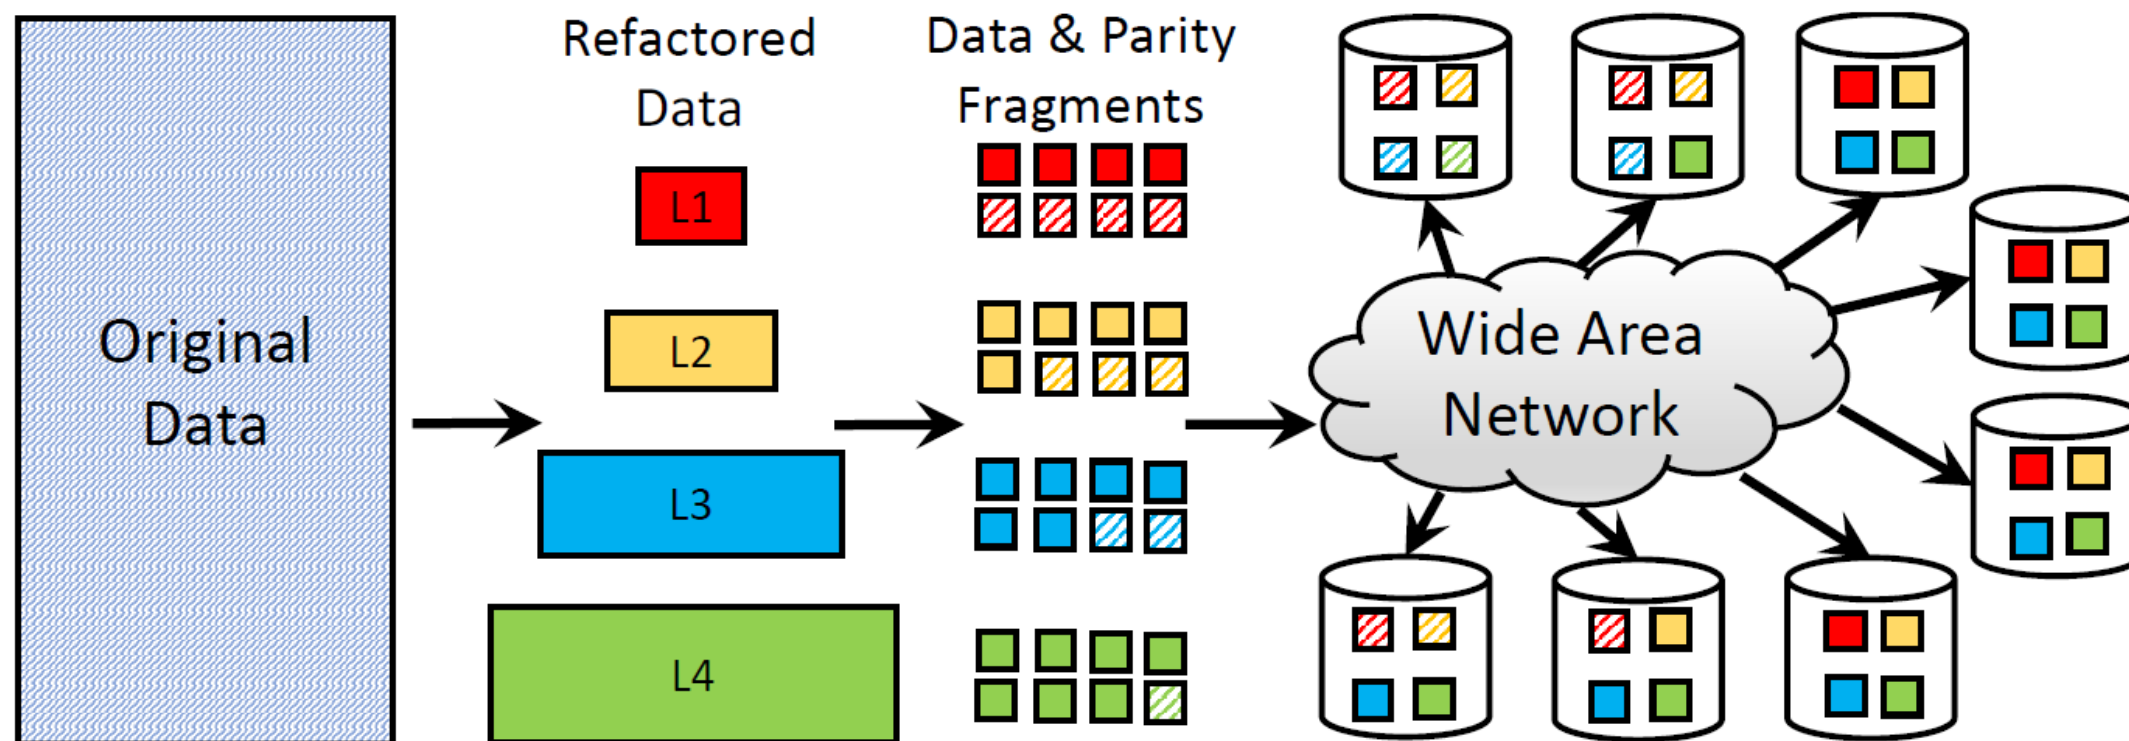
**Still, it is expensive to create, transfer, and store many (large) parity and data fragments to achieve high availability for large scientific datasets.**

## Erasure Coding

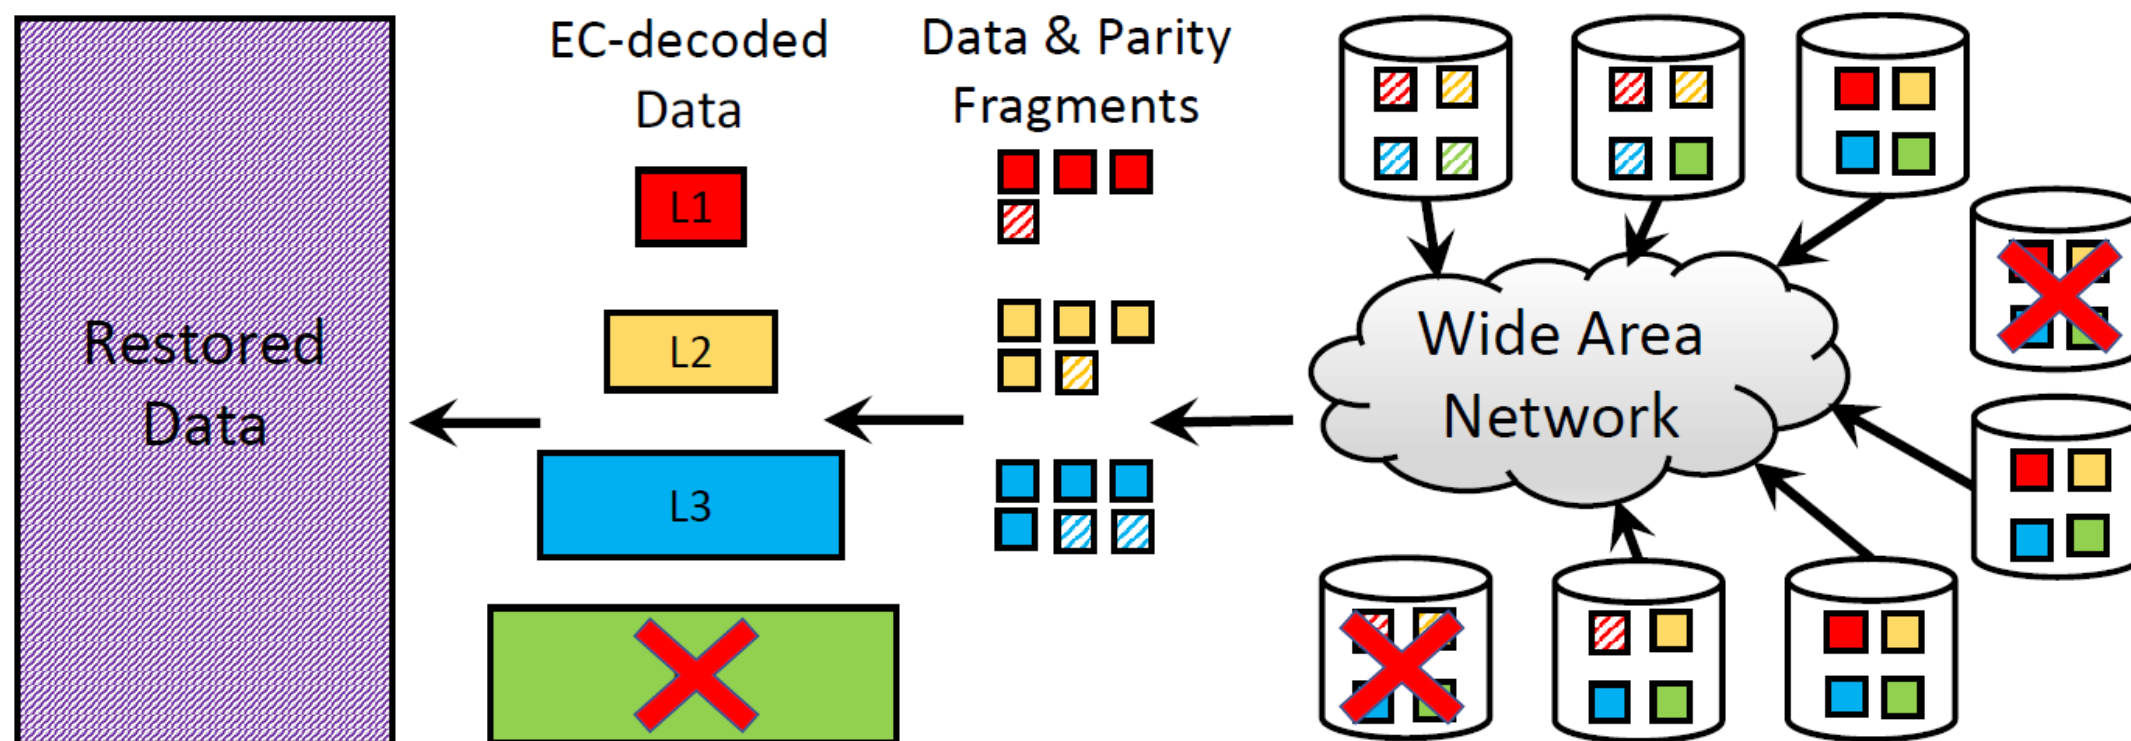
- Splitting data into data and parity fragments
- Data restoration from a subset of fragments



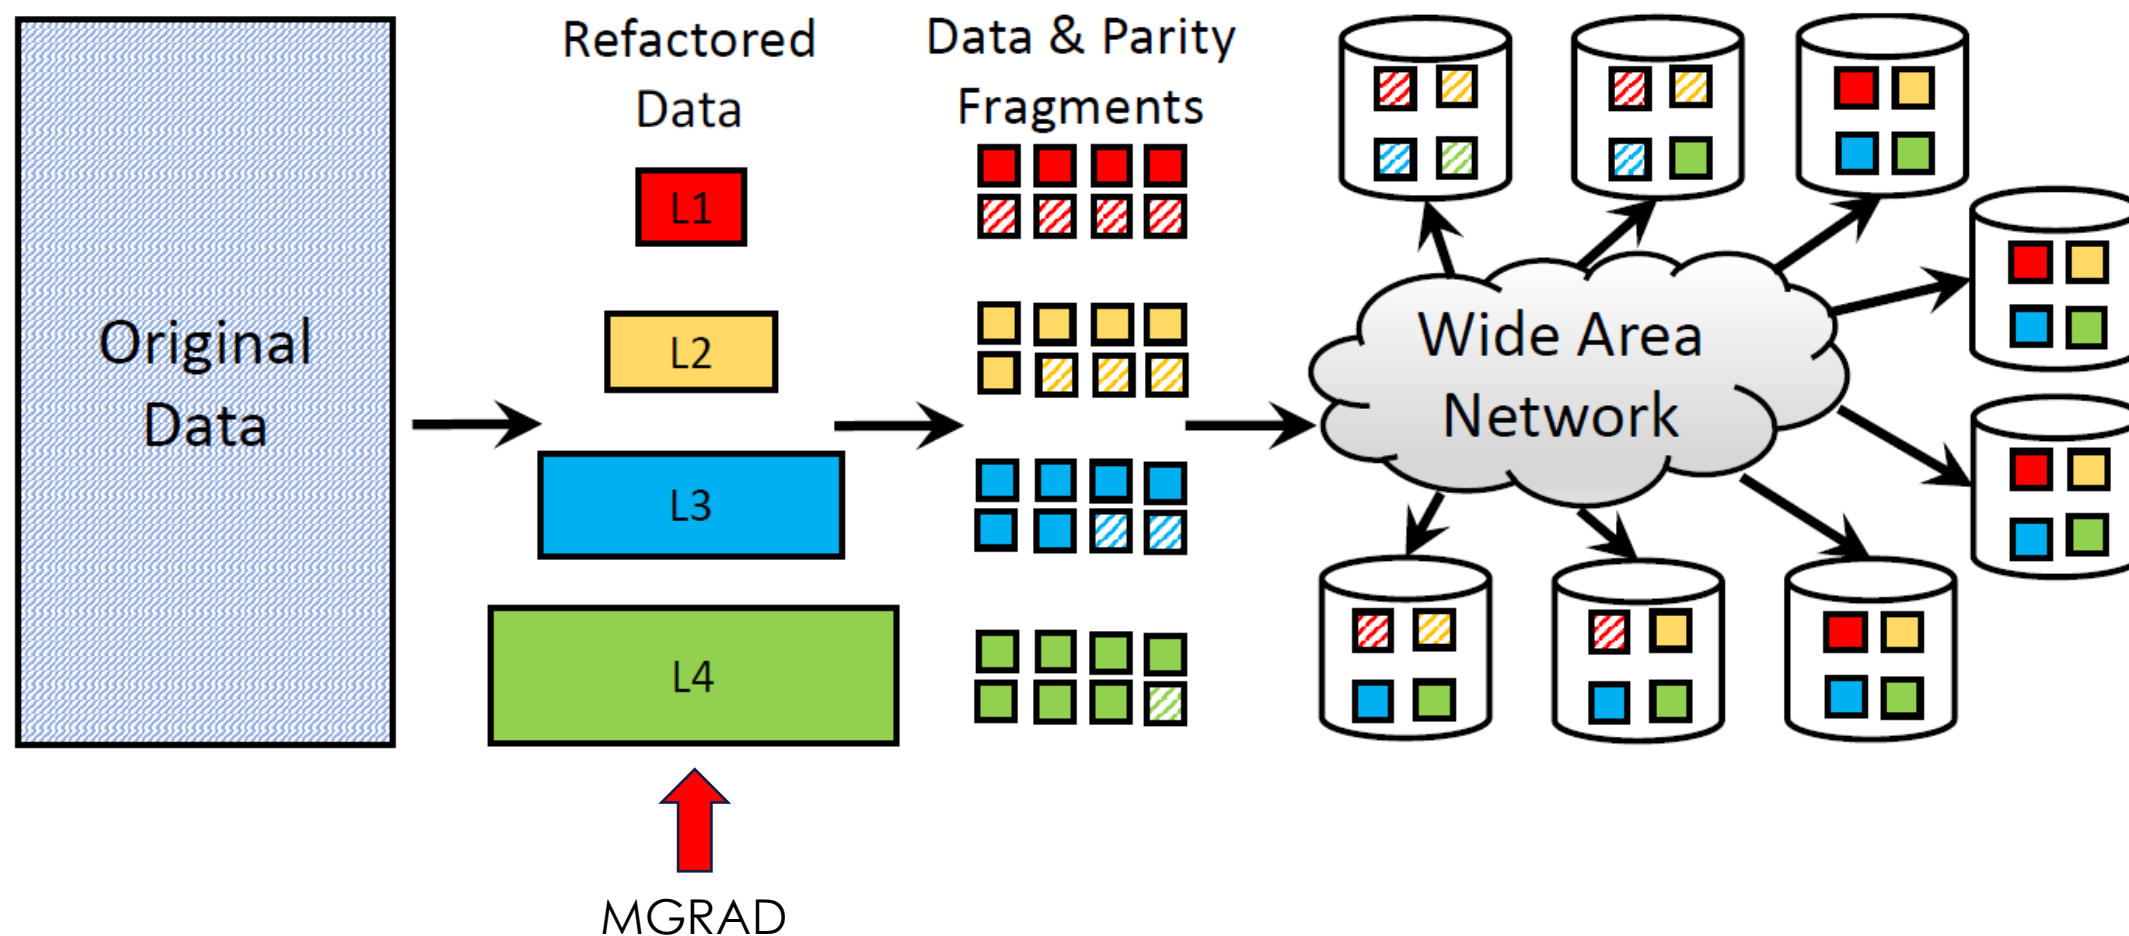
# RAPIDS – Store Data



# RAPIDS – Get Data



# Refactor & Compress



# Refactor & Compress

- Error-bounded lossy compression
- Refactor the original data into a hierarchical representation
  - E.g., 4 levels
  - Size:  $s_1 < s_2 < s_3 < s_4$
- Data reduction is applied
  - Error in the reduced representation
  - Size:  $S > s_1 + s_2 + s_3 + s_4$

Refactored  
Data

L1

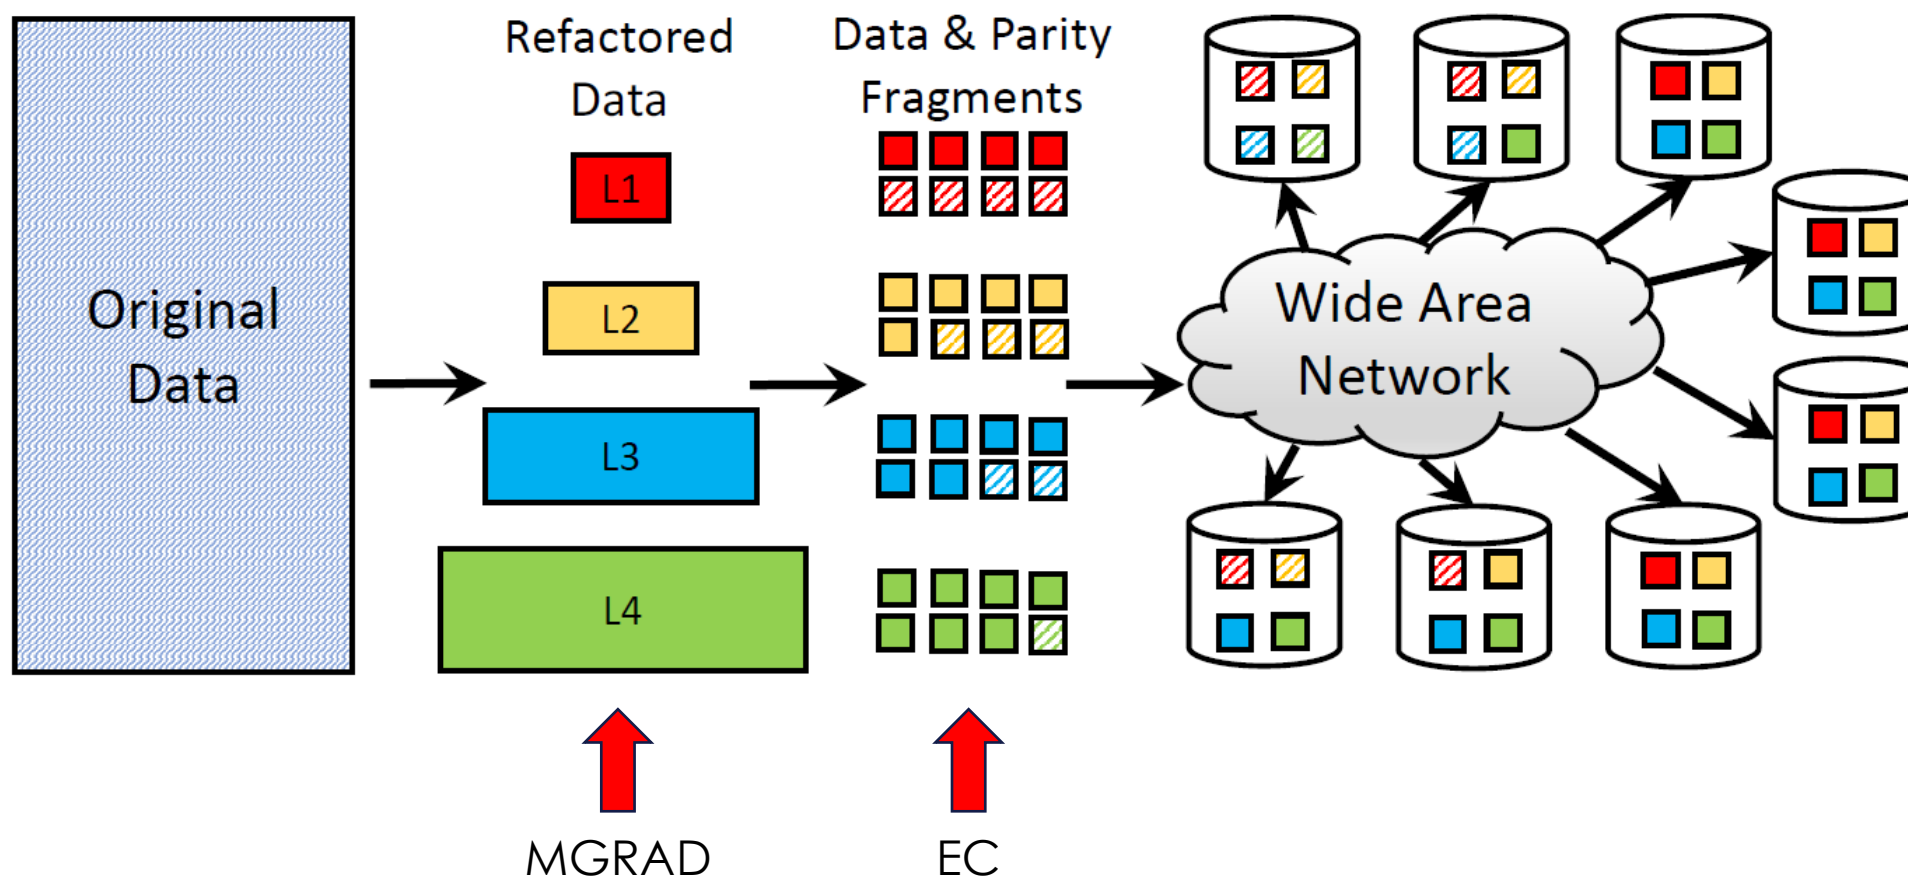
L2

L3

L4





# Erasure Coding



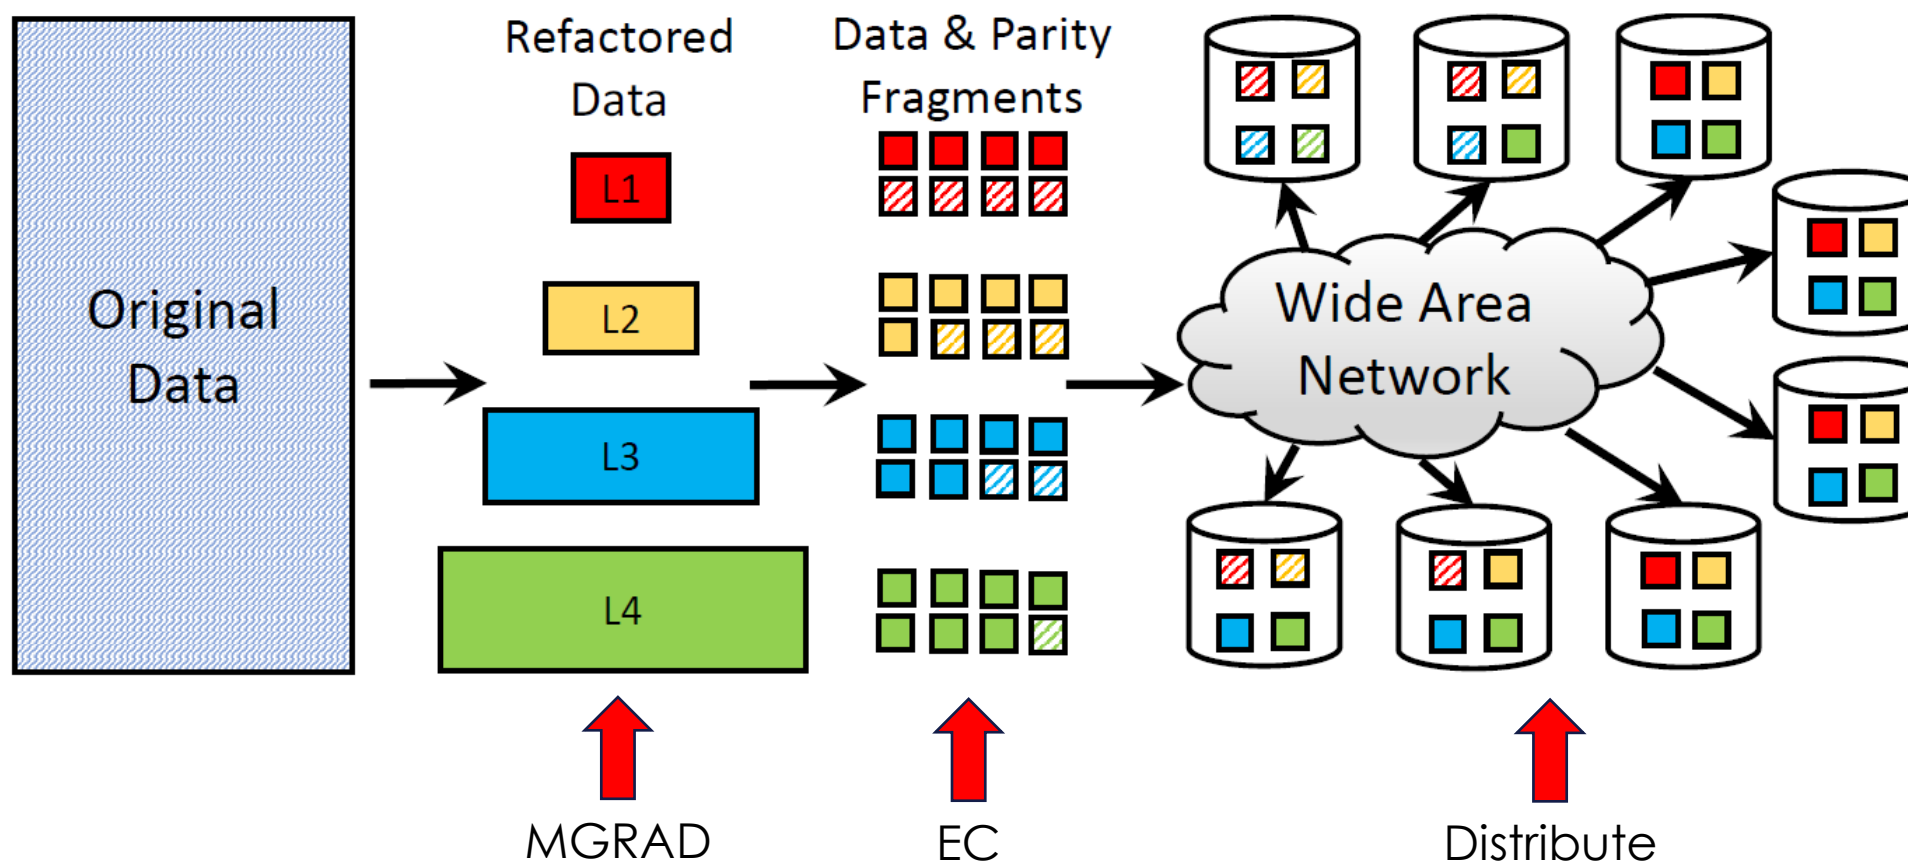
# Erasure Coding

- Erasure coding (each level) with fault tolerance
  - Encoded into  $n - m_j$  data and  $m_j$  parity
  - Parity  $m_1 > m_2 > m_3 > m_4$
- Reduces storage and network overhead compared to whole data compression
  - Data  $s_1 < s_2 < s_3 < s_4$ 
    - Smaller duplication overhead.

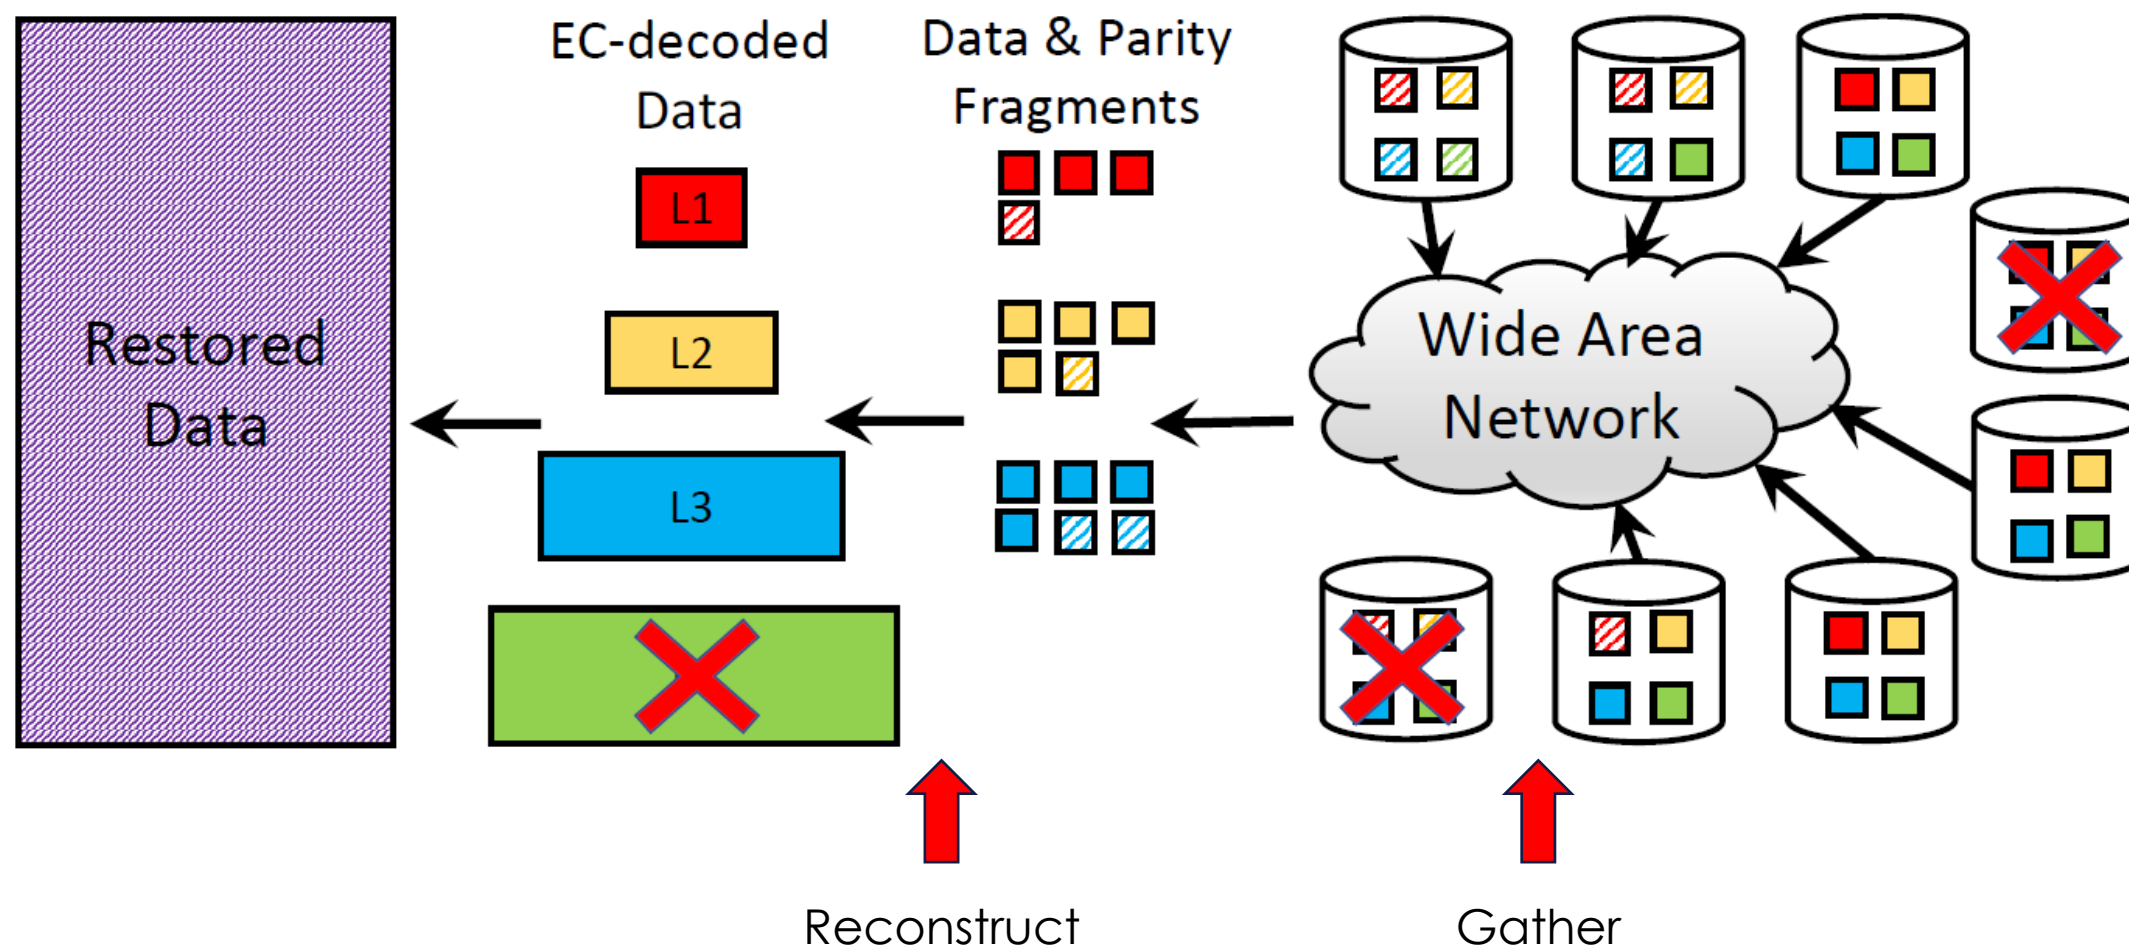
## Data & Parity Fragments



# Distributed Storage



# RAPIDS – Get Data




# RAPIDS – Get Data

- Data gathering optimization
  - Which fragment fetched from which source
    - Network overhead / latency
- Reconstruction – Fault tolerance configuration is critical for availability and accuracy
  - Minimize the expected relative error

# Data Gathering

- Find a set with values of enough fragments (fragments, systems) to minimize data transfer time
  - Find time for each transfer request (including multiple request sources)
  - Reduce average time for all requests
- Fetch a subset of fragments
  - Prioritize high-bandwidth
  - Reduce concurrency



# Fault Tolerance

- Erasure coding defines data restructuring
  - how many failures each level can sustain
- If failure  $(N) > m_1$ ,
  - Data will be unavailable for reconstruction
  - Max error (1), data is useless
- If  $m_{j+1} < N \leq m_j$ ,
  - Data will be available for reconstruction
  - Error  $e_j$

## Data & Parity

### Fragments



# Fault Tolerance

- Minimize faults (each level) to reduce error under constraints (optimization)
  - Storage overhead should not exceed the threshold (defined)
  - Number of failures at each level can tolerate should follow  $m_1 > m_2 > \dots > m_l \geq 1$  as  $(s_1 < s_2 < \dots < s_l)$
- Small solution space – Brute-force search
  - $O(\text{Storage systems} - \text{refactor levels})^4$
- Large solution space – Heuristic search



# Fault Tolerance Heuristic



1. Set initial fault tolerance configuration w.r.t. overall storage overhead – reducing search space
2. Increment parity on the level (bottom start)
3. If storage overhead is violated
  - i. Reduce/use the last stable configuration
  - ii. Move to the level above in the hierarchy and go to step 2

## Assumptions

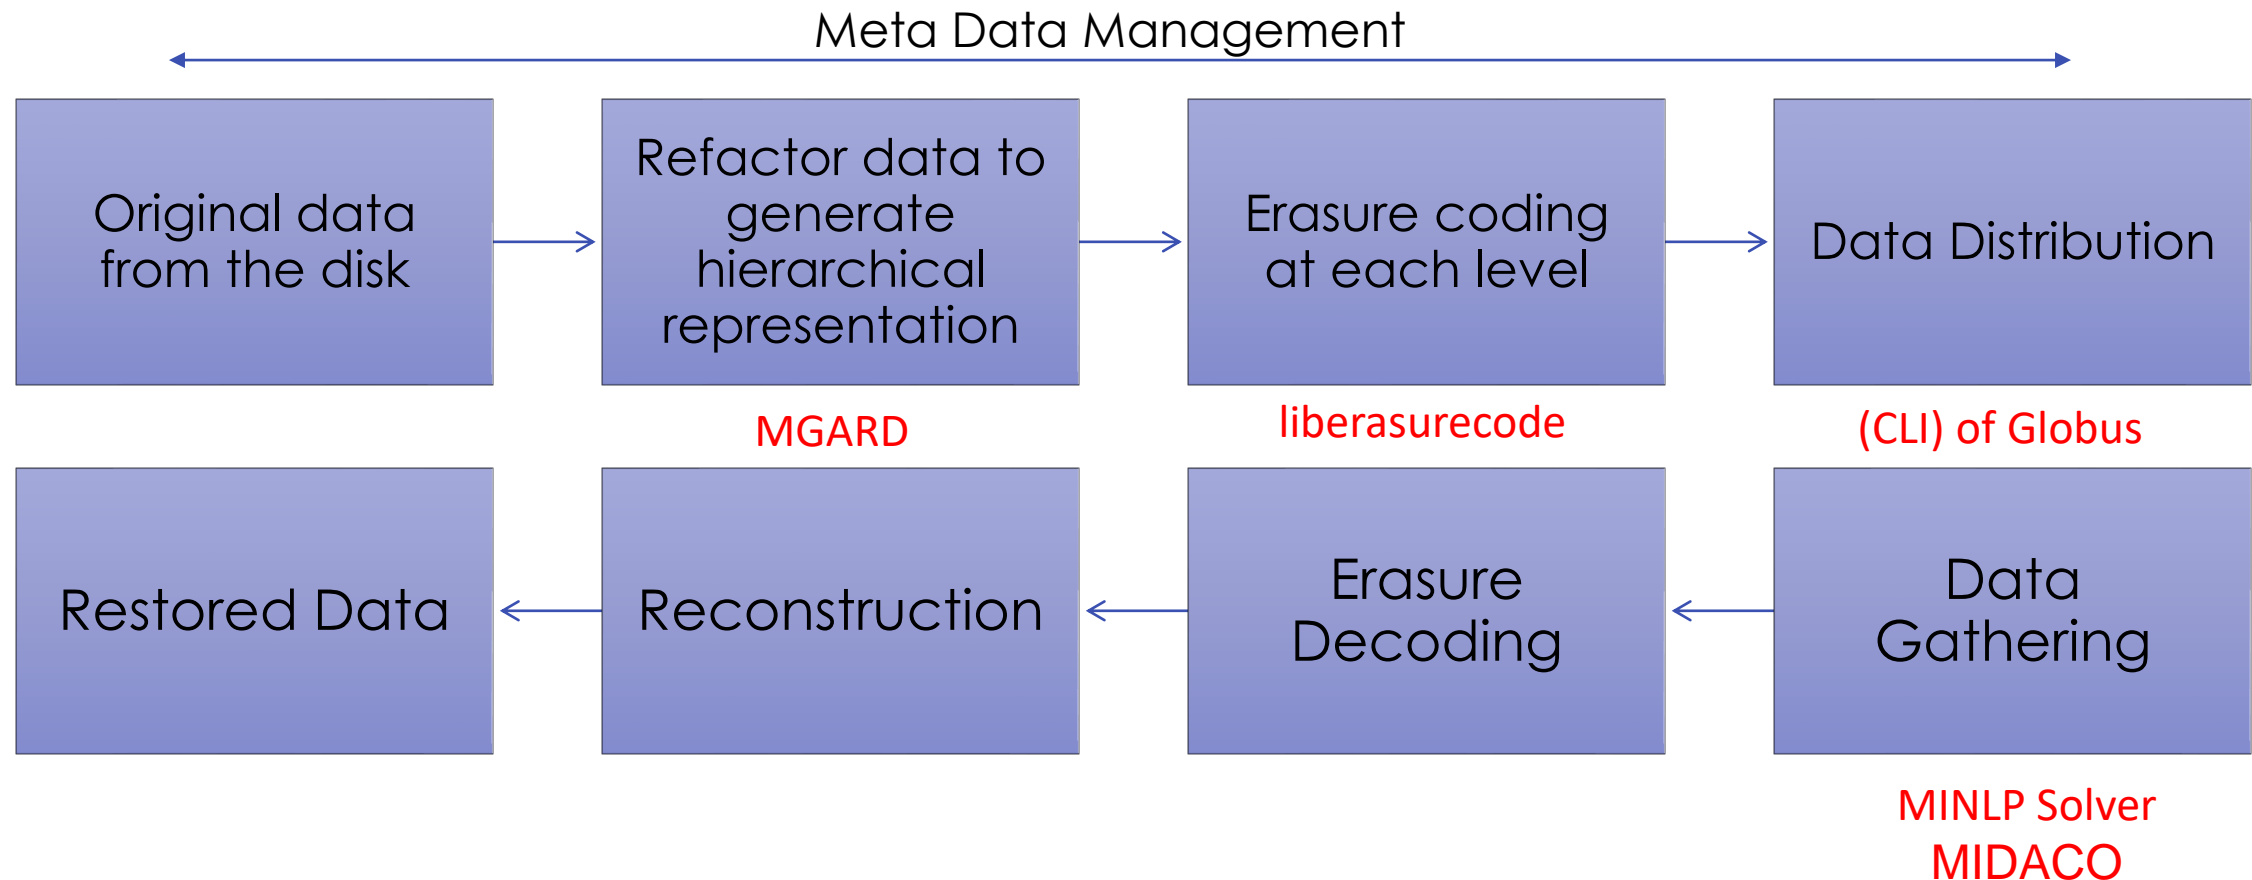
- i.  $s1 \ll s2 \ll s3 \ll s4$
- ii.  $e1 \gg e2 \gg e3 \gg e4$

## Data & Parity

### Fragments



# Recap



# Meta Data Management

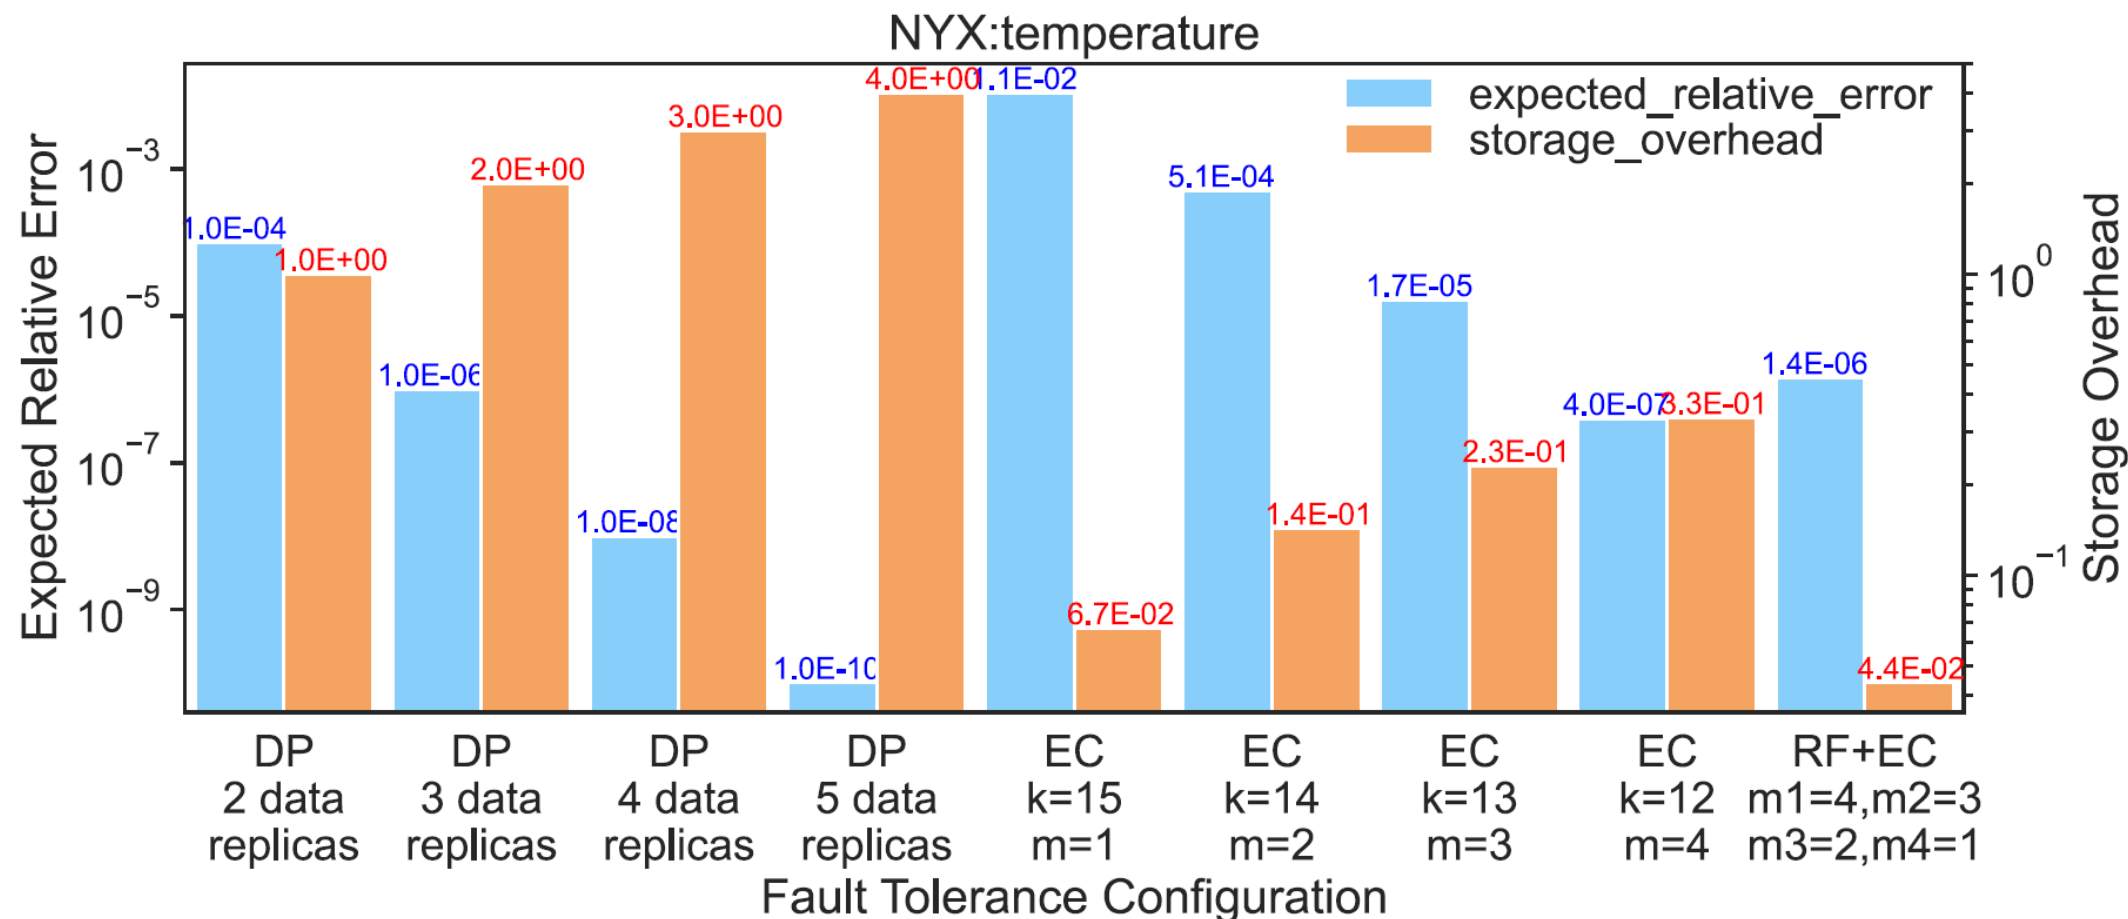
- Track all the data and parity fragments
  - key-value database (fragment: system)
  - Restoring lost fragments – EC (fragment: system)
  - Reconstruction information
  - Central – Prone to failure

# Evaluation – Data Sets


- NYX cosmology simulation
- SCALE-LETKF weather simulation
- Hurricane Isabel climate simulation
- Oak Ridge Leadership Computing Facility (OLCF)
  - 704-compute node cluster; where each node has
    - 2 16-core Processors (32 cores)
    - 256GB of memory
  - 9 GPU nodes; where each node has
    - 2 NVIDIA K80 GPUs, 2 Intel Xeon 14-core processors
    - 1TB memory

Dataset	Object name	Size/object
NYX	temperature	16TB
NYX	velocity_x	16TB
SCALE-LETKF	PRES	16.82TB
SCALE-LETKF	T	16.82TB
Hurricane Isabel	Pf48.bin	2.98TB
Hurricane Isabel	TCf48.bin	2.98TB

# Evaluation – Error and Storage




# Evaluation – Latency



# Evaluation – Data Gathering

Average over 50 iterations



Random: varying the seed

Naïve: Sort by bandwidth and select from top-bottom


Optimized: Naïve + MIDACO

# Evaluation – FT Heuristic


Data object	Optimal FT Configuration		Speedup ( $t_{BF}/t_H$ )
	Brute-Force	Heuristic	
NYX:temperature	[8,5,4,2]	[8,5,4,2]	124
NYX:velocity_x	[5,4,3,1]	[5,4,3,1]	152
SCALE:PRES	[9,6,4,2]	[9,6,4,2]	114
SCALE:T	[7,4,3,2]	[7,4,3,2]	137
hurricane:Pf48.bin	[11,10,8,7]	[11,10,8,7]	152
hurricane:TCf48.bin	[11,9,8,6]	[11,9,8,6]	137



# Evaluation – Performance



# Evaluation – Performance



# RAPIDS

- Two optimization models.
  - A mixed integer nonlinear programming (MINLP) solver for gathering data – MIDACO
  - A heuristic algorithm compared to the brute-force search for optimizing fault tolerance
- Evaluation results show
  - Minimum expected relative error
  - Reduced storage overhead
  - Higher data latency

Thank You

# Optimal Fault Tolerance Configurations

**Input:** Initial fault tolerance configurations  $M^i = [m_1^i, m_2^i, \dots, m_l^i]$

**Output:** Optimal fault tolerance configurations  $M^o = [m_1^o, m_2^o, \dots, m_l^o]$

$M = [m_1, m_2, \dots, m_l] = M^i, l_{curr} = l, M_{prev} = [];$

**while** *True* **do**

$$W = \frac{\sum_{j=1}^l \frac{m_j}{n-m_j} s_j}{S};$$

**if**  $W < \omega$  **then**

**foreach**  $1 \leq x < l_{curr}$  **do**

$m_x = m_x + 1;$

**end**

$M = [m_1, m_2, \dots, m_l];$

**else**

**foreach**  $1 \leq x < l_{curr}$  **do**

$m_x = m_x - 1;$

**end**

$M = [m_1, m_2, \dots, m_l];$

$l_{curr} = l_{curr} - 1;$

**end**

**if**  $M == M_{prev}$  **then**

**break;**

**end**

$M_{prev} = M;$

**end**

$M^o = M;$

# PMGRAD

- PMGARD, multilevel decomposition approach, with an error-controlled data refactoring and reconstruction framework for scientific data.
  - Decomposing original data into multiple components
  - Multilevel coefficients, at relative importance to the precision of the reconstructed data, are used to reconstruct the approximation of the original data